

Chapter 4: Threads & Concurrency

these slides are edited and some of the contents are taken from
<https://www.scs.stanford.edu/24wi-cs212/notes/processes.pdf>
and
<http://www.it.uu.se/education/course/homepage/os/vt18/module-4/implementing-threads/>

Main two point:

- 1) threads vs. process
- 2) kernel-level vs. user-level threads

Review of process


Concurrency and threads

- Overview
- Multicore Programming

Review: Processes

- A process is an instance of a program running
- Why processes?
 - Simplicity of programming
 - Speed: Higher throughput, lower latency

Speed



Multiple processes can increase **CPU utilization**

Multiple processes can reduce **latency**



Running A then B requires 100 sec for B to complete



A is slower than if it had whole machine to itself,
but still ≤ 100 sec unless both A and B completely CPU-bound

Multitasking in real world

- 1 worker 10 months to make 1 widget
- hire 100 workers to make 100 widgets
 - Latency for first widget $\gg 1/10$ month
 - setup/learning etc
 - Throughput may be < 10 widgets per month
 - **if can't perfectly parallelize task**
- Or 100 workers making 10,000 widgets may achieve > 10 widgets/month
 - **if workers never idly wait**
 - **capitalism at its best** 🙄

A **process's view** of the world

Each process has own view of machine

- Its own address space – `*(char *)0xc000` different in P1 & P2
- Its own open files
- Its own virtual CPU (through preemptive
- multitasking)

Simplifies programming model

- gcc does not care that firefox is running

Sometimes want interaction between processes

- Simplest is through files: emacs edits file, gcc compiles it
- More complicated: Shell/command, Window manager/app.




Figure 3.1 Layout of a process in memory.


Kernel's view of processes: implementing a process

- Keep a data structure for each process
 - Process Control Block (PCB)
 - Called `proc` in Unix, `task_struct` in Linux,
- Tracks state of the process
 - Running, ready (runnable), waiting, etc.
- Includes information necessary to run
 - Registers, virtual memory mappings, etc.
 - Open files (including memory mapped files)
- Various other data about the process
 - Credentials (user/group ID), signal mask, controlling terminal, priority, accounting statistics, whether being debugged, which system call binary emulation in use, . . .

Process state
Process ID
User id, etc.
Program counter
Registers
Address space (VM data structs)
Open files

PCB

Process states



Process can be in one of several states

- new & terminated at beginning & end of life
- running – currently executing (or will execute on kernel return)
- ready – can run, but kernel has chosen different process to run
- waiting – needs async event (e.g., disk operation) to proceed

Which process should kernel run?

- if 0 runnable, run idle loop (or halt CPU), if 1 runnable, run it
- if >1 runnable, must make scheduling decision

Preemption

Can preempt a process **when kernel gets control**

- **Running process can vector control to kernel**

- System call, page fault, illegal instruction, etc.
- May put current process to sleep
 - e.g., read from disk
- May make other process runnable
 - e.g., fork, write to pipe

- **Periodic timer interrupt**

- If running process used up quantum, schedule another

- **Device interrupt**


- Disk request completed, or packet arrived on network
- Previously waiting process becomes runnable
- Schedule if higher priority than current running proc.

Changing running process is called a **context switch**

Context switch

Very machine dependent. Typical things include:


- Save program counter and integer registers (always)
- Save floating point or other special registers
- Save condition codes
- Change virtual address translations




Context switch: Cost

Non-negligible cost

- Save/restore floating point registers expensive
 - Optimization: only save if process used floating point
- May require **flushing TLB** (memory translation hardware)
- Usually causes more cache misses (switch working sets)



Inter-process communication



- How can processes interact in real time?
 - (a) By passing messages through the kernel
 - (b) By sharing a region of physical memory
 - (c) Through asynchronous signals or alerts

Creating/deleting processes in Unix

- `int fork (void);`
 - Create new process that is exact copy of current one
 - Returns *process ID* of new process in “parent”
 - Returns 0 in “child”
- `int waitpid (int pid, int *stat, int opt);`
 - `pid` – process to wait for, or -1 for any
 - `stat` – will contain exit value, or signal
 - `opt` – usually 0 or `WNOHANG`
 - Returns process ID or -1 on error
- `void exit (int status);`
 - Current process ceases to exist
 - `status` shows up in `waitpid` (shifted)
 - By convention, `status` of 0 is success, non-zero error
- `int kill (int pid, int sig);`
 - Sends signal `sig` to process `pid`
 - `SIGTERM` most common value, kills process by default (but application can catch it for “cleanup”)
 - `SIGKILL` stronger, kills process always

Other examples

E.g. windows system call [CreateProcessA function \(processthreadsapi.h\) - Win32 apps | Microsoft Learn](#)

[CreateProcessAsUserA function \(processthreadsapi.h\) - Win32 apps | Microsoft Learn](#)

```
BOOL CreateProcessAsUserA(  
    [in, optional]    HANDLE          hToken,  
    [in, optional]    LPCSTR          lpApplicationName,  
    [in, out, optional] LPSTR          lpCommandLine,  
    [in, optional]    LPSECURITY_ATTRIBUTES lpProcessAttributes,  
    [in, optional]    LPSECURITY_ATTRIBUTES lpThreadAttributes,  
    [in]              BOOL             bInheritHandles,  
    [in]              DWORD            dwCreationFlags,  
    [in, optional]    LPVOID           lpEnvironment,  
    [in, optional]    LPCSTR           lpCurrentDirectory,  
    [in]              LPSTARTUPINFOA   lpStartupInfo,  
    [out]             LPPROCESS_INFORMATION lpProcessInformation  
);
```

Running programs

- `int exece (char *prog, char **argv, char **envp);`
 - `prog` – full pathname of program to run
 - `argv` – argument vector that gets passed to `main`
 - `envp` – environment variables, e.g., `PATH`, `HOME`
- **Generally called through a wrapper functions**

- `int execvp (char *prog, char **argv);`
Search `PATH` for `prog`, use current environment
- `int execlp (char *prog, char *arg, ...);`
List arguments one at a time, finish with `NULL`

```
pid_t pid; char **av;  
void doexec () {  
    execvp (av[0], av);  
    perror (av[0]);  
    exit (1);  
}
```

```
/* ... main loop: */  
for (;;) {  
    parse_next_line_of_input (&av, stdin);  
    switch (pid = fork ()) {  
        case -1:  
            perror ("fork"); break;  
        case 0:  
            doexec ();  
        default:  
            waitpid (pid, NULL, 0); break;  
    }  
}
```

- `int dup2 (int oldfd, int newfd);`
 - Closes `newfd`, if it was a valid descriptor
 - Makes `newfd` an exact copy of `oldfd`
 - Two file descriptors will share same offset (lseek on one will affect both)
- `int fcntl (int fd, int cmd, ...)` – **misc fd configuration**
 - `fcntl (fd, F_SETFD, val)` – sets close-on-exec flag
When `val == 0`, fd not inherited by spawned programs
 - `fcntl (fd, F_GETFL)` – get misc fd flags
 - `fcntl (fd, F_SETFL, val)` – set misc fd flags

Loop that reads a command and executes it

Recognizes

`command < input > output 2> errlog`

```
void doexec (void) {
    int fd;
    if (infile) {          /* non-NULL for "command < infile" */
        if ((fd = open (infile, O_RDONLY)) < 0) {
            perror (infile);
            exit (1);
        }
        if (fd != 0) {
            dup2 (fd, 0);
            close (fd);
        }
    }

    /* ... do same for outfile→fd 1, errfile→fd 2 ... */

    execvp (av[0], av);
    perror (av[0]);
    exit (1);
}
```

Example IPC: Pipes-message passing through kernel

- `int pipe (int fds[2]);`
 - Returns two file descriptors in `fds[0]` and `fds[1]`
 - Data written to `fds[1]` will be returned by `read` on `fds[0]`
 - When last copy of `fds[1]` closed, `fds[0]` will return EOF
 - Returns 0 on success, -1 on error

- **Operations on pipes**

- `read/write/close` – as with files
- When `fds[1]` closed, `read(fds[0])` returns 0 bytes
- When `fds[0]` closed, `write(fds[1])`:
 - ▷ Kills process with `SIGPIPE`
 - ▷ Or if signal ignored, fails with `EPIPE`

```
void doexec (void) {
    while (outcmd) {
        int pipefds[2]; pipe (pipefds);
        switch (fork ()) {
            case -1:
                perror ("fork"); exit (1);
            case 0:
                dup2 (pipefds[1], 1);
                close (pipefds[0]); close (pipefds[1]);
                outcmd = NULL;
                break;
            default:
                dup2 (pipefds[0], 0);
                close (pipefds[0]); close (pipefds[1]);
                parse_command_line (&av, &outcmd, outcmd);
                break;
        }
    }
}
```


Overview of multicore programming

- Overview
- Multicore Programming
- Multithreading Models

Multicore Programming


- **Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:
 - **Dividing activities**
 - **Balance**
 - **Data splitting**
 - **Data dependency**
 - **Testing and debugging**

Multicore Programming: key definitions


- **Parallelism** implies a system can perform more than one task simultaneously
- Operations are **concurrent** if they can progress independently at the same time
- **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency

Concurrency (think as logical) vs. Parallelism (actual)

- **Concurrent execution on single-core system:**



- **Parallelism on a multi-core system:**




Serial vs Parallel computing

A process is working on multiple tasks at the same time


Parallelism: splitting tasks into subtasks

Parallel computing: running subtasks
(concurrent parts) parallel in multiple processor

Serial computing




Parallel computing




Multicore Programming: Types of parallelism

- **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each



- **Task parallelism** – distributes threads across cores, each thread performing unique operation



Matrix addition: Concurrent parts?

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}$$

$$\mathbf{B} = \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1n} \\ b_{21} & b_{22} & \cdots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{m1} & b_{m2} & \cdots & b_{mn} \end{bmatrix}$$

$$\mathbf{A} + \mathbf{B} = \begin{bmatrix} a_{11} + b_{11} & a_{12} + b_{12} & \cdots & a_{1n} + b_{1n} \\ a_{21} + b_{21} & a_{22} + b_{22} & \cdots & a_{2n} + b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} + b_{m1} & a_{m2} + b_{m2} & \cdots & a_{mn} + b_{mn} \end{bmatrix}$$

Data: $\mathbf{a}_{ij}, \mathbf{b}_{ij}$

Task: Addition

The same task, different data!

Matrix multiplication: Concurrent parts?

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}$$

$$\mathbf{B} = \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1n} \\ b_{21} & b_{22} & \cdots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{m1} & b_{m2} & \cdots & b_{mn} \end{bmatrix}$$

$$\mathbf{C} = \mathbf{AB}$$

$$\mathbf{C} = \begin{pmatrix} a_{11}b_{11} + \cdots + a_{1n}b_{n1} & a_{11}b_{12} + \cdots + a_{1n}b_{n2} & \cdots & a_{11}b_{1p} + \cdots + a_{1n}b_{np} \\ a_{21}b_{11} + \cdots + a_{2n}b_{n1} & a_{21}b_{12} + \cdots + a_{2n}b_{n2} & \cdots & a_{21}b_{1p} + \cdots + a_{2n}b_{np} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1}b_{11} + \cdots + a_{mn}b_{n1} & a_{m1}b_{12} + \cdots + a_{mn}b_{n2} & \cdots & a_{m1}b_{1p} + \cdots + a_{mn}b_{np} \end{pmatrix}$$

Data: **a_i**, **B** or **A**, **b_j**

Task: Multiplication(and addition)

The same task, different data!

What if we want to do both?

Matrix addition

$$\mathbf{T} = \mathbf{A} + \mathbf{B}$$

Matrix multiplication


$$\mathbf{C} = \mathbf{A}\mathbf{B}$$

Data: $\mathbf{A}$, $\mathbf{B}$

Task: Addition, Multiplication


The same data, different concurrent tasks!

Speeding-up calculations: Moore's law



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2021 by K. Rupp

How much can we speed-up?



[Amdahl's law - Wikipedia](#)

Amdahl's Law in more details...


- Identifies performance gains from adding additional cores to an application that has both serial and parallel components

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$


- S is fraction of task that is necessarily serial (the rest is parallel)
- N processing cores

- What is the speedup, if application is 75% parallel and 25% serial, moving from 1 to 2 cores?
- What happens
 - as S approaches 0?
 - as S approaches 1?
 - as N approaches infinity?

How to implement threads?



single-threaded process



multithreaded process

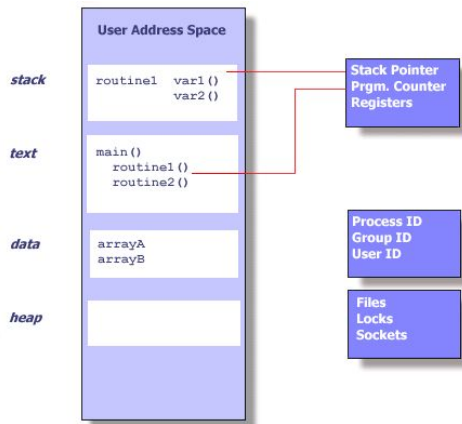
- A thread is a schedulable execution context
 - Program counter, stack, registers, . . .
- Simple programs use one thread per process
- But can also have multi-threaded programs
 - Multiple threads running in same process's address space

Why need threads: Processes vs Threads


Thread maintains only its own

Process

- Each process has its own virtual address space
- **fork() is expensive**




- Stack pointer
- Registers
- Scheduling properties (such as policy or priority)
- Set of pending and blocked signals
- Thread specific data.



More Motivation for threads

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Allows one process to use multiple CPUs or cores
 - Allows program to overlap I/O and computation
- Process creation is heavy-weight while thread creation is light-weight
 - All threads in one process share memory, file descriptors, etc.
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

Example: Multithreaded Server Architecture



```
for (;;) {  
    fd = accept_client ();  
    thread_create (service_client, &fd);  
}
```


Summary of Threading-Benefits

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- **Scalability** – process can take advantage of multicore architectures